

Attempt to Reconstruct Confidential Data Using Public Data

Sarah Birmingham

Coding It Forward Data Science Fellow

Census of Fatal Occupational Injuries

August 5th 2021

Supervisor:

Ellen Galantucci, PhD

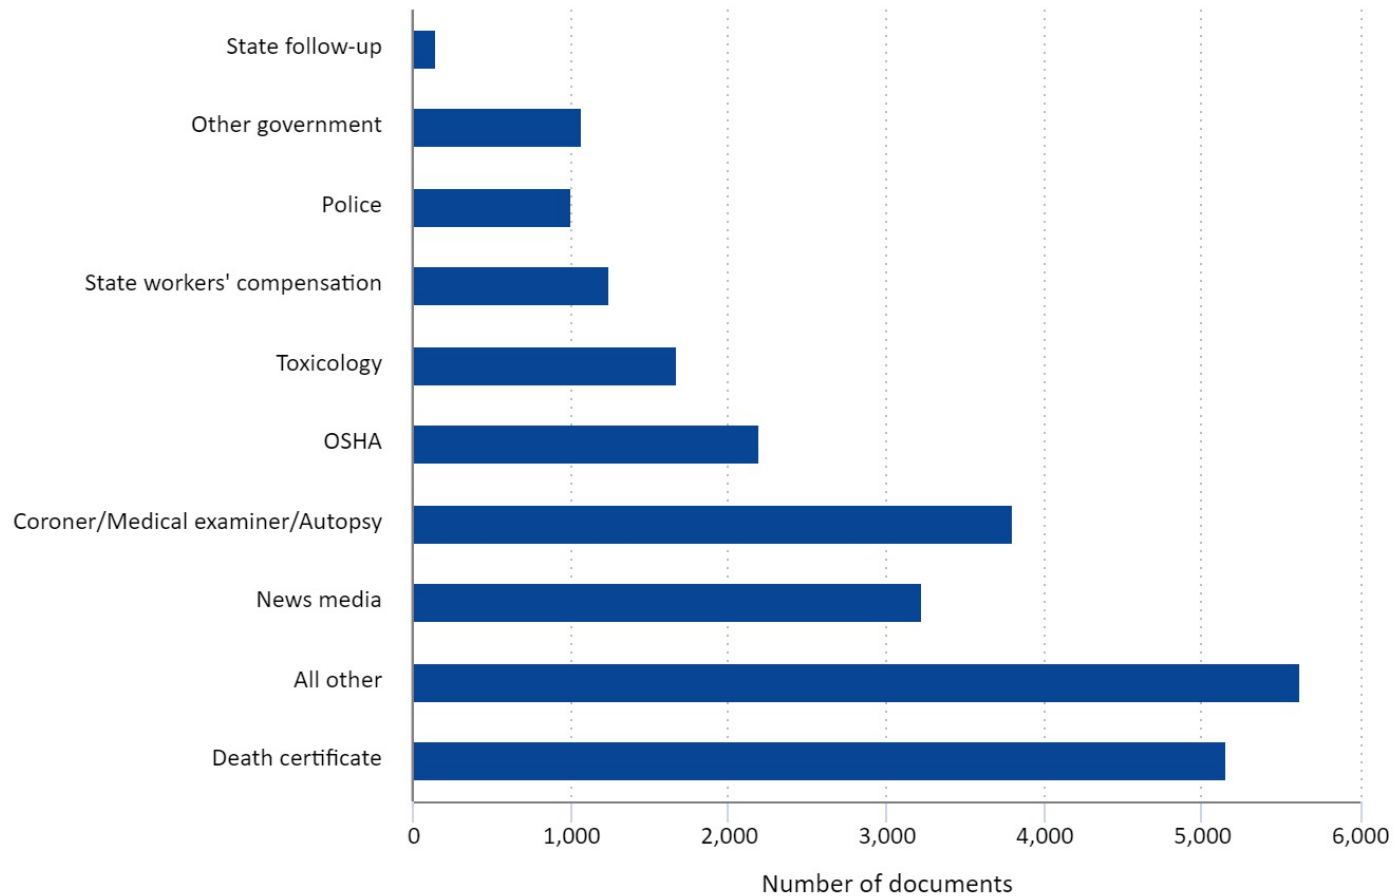
Mathematical Statistician



What is CFOI?

- The Census of Fatal Occupational Injuries is a count of fatal work injuries in the United States and measures the number of people who die at work or due to a work-related injury
- CFOI data is used by OSHA and NIOSH for regulatory and advisory purposes
- Data is obtained from a variety of sources and BLS is obligated to keep identifiable information confidential under CIPSEA

Sources of data on fatal work injuries, 2019



Hover over chart to view data.

Source: U.S. Bureau of Labor Statistics.

What are database reconstruction attacks (DRA)?

- Published data is used to infer underlying microdata
- Researchers at UT-Austin used the Netflix Prize dataset and public information from IMDB to identify Netflix records of known users

- CFOI publishes decedent data on:
 - ▶ Type, occupation, industry, and sector of employment
 - ▶ Gender, age, and race and Hispanic origin status
 - ▶ Causal event and state of occurrence
- Data is available on a national and per-state basis
- Need to balance usefulness of published data with responsibility to protect confidentiality
- Publishability requirements were changed in 2019
- Data is hard to protect because counts are small

Problem Set-Up

- Access to both public and confidential microdata;
able to start with all known public data
 - ▶ Unrealistic for an actual adversarial attack
- Extract assignment rules from published tabular data
 - ▶ Rule examples include:
 - Four gender_2 in occupation_53
 - No government employees were self-employed
 - If not gender_1, gender_2

Record ID	State	Event	Occupation	Industry	Public/private	Wage/self-employed	Gender	Age category	Race
1	42	1	47	238	50	2	1	4	1
2	42	2	53	926	10	4	1	4	2
3	42	5	47	485	50	4	1	5	6
4	42	4	11	238	50	2	2	6	3
5	42	3	53	561	50	4	1	3	1
6	42	3	49	713	50	4	2	4	1
7	42	6	31	621	50	4	1	4	2

Initial Approach: Backtracking

- Conceptualized the problem as an enormous Sudoku puzzle or the game “Clue”: starting from known information, make initial inferences and eliminate other possibilities
- Assign values until I reach a solution or a dead end; in which case, backtrack

Record ID	State	Event	Occupation	Industry	Public/private	Wage/self-employed	Gender	Age category	Race
1	42	1	11	238	50	2	1	5	3
2	42	2			10				
3	42	5			50			5	6
4	42	4		238	50			6	
5	42	3							
6	42			713			2		
7	42		31						

Rule 1: Only one race_3

Rule 2: Occupation_11 has race_3

Rule 3: Race_3 has gender_2

Did it work?

No

- Problem size was way too big for this to be a feasible approach, even on individual states
- Sudoku on a board with $> 50,000$ cells
- Tried to modify this attempt by:
 - ▶ Assigning values pre-approved pairs
 - ▶ Reordering columns to fill in easier-to-guess cells first



Second Approach:

Optimization with OR-Tools library

- Google OR-Tools is an open-source software suite for optimization problems
- Create a model by assigning variables and constraints
- Use a SAT solver (built-in, open-source, or commercial) to solve

Record ID	State	Event_1	Event_2	Event_3	Event_4	Event_5	Event_6	Event_?
1	42	(0,1)	(0,1)	(0,1)	(0,1)	(0,1)	(0,1)	(0,1)
2	42	0	1	0	0	0	0	0
3	42	0	0	0	0	1	0	0
4	42	0	0	0	1	0	0	0
5	42	0	0	1	0	0	0	0
6	42	(0,1)	(0,1)	(0,1)	(0,1)	(0,1)	(0,1)	(0,1)
7	42	(0,1)	(0,1)	(0,1)	(0,1)	(0,1)	(0,1)	(0,1)

Did it work?

No

- Found one potential solution; couldn't find optimal or multiple solutions
- Didn't seem designed to handle problems that had an inconsistent number of assignable values
- No counting function



Third Approach: Merging Permutations

- Use tabular data to find # and value of assignable options
- Necessarily have some unknown values-will never be able to solve for those
- Create a list of valid permutations for each category, keeping public values fixed



Third Approach: Merging Permutations

- For each grouping of two categories $[I,J]$ for which a 'rule' exists, eliminate members of both categories that don't satisfy that rule
 - ▶ No rule between categories \rightarrow all pairs are valid
- For each grouping of two pairs $[I,J],[J,K]$ that share a common category, eliminate members that don't:
 - ▶ Satisfy any rules that exist between all three $[i,j,k]$
 - ▶ Have $[i,k]$ in list of pairs $[I,K]$

	Event	Industry	Public/ private	Wage/ self- employed	
	3	238	50	2	
	2	985	10	4	
	5	985	50	4	
	4	238	50	4	
	3	565	50	4	
	5	713	50	4	
	6	621	50	4	
n perm	6	23	1	5	Total
					218

Did it work?

Maybe!

- Continue to merge until you have a group of all nine categories
 - ▶ Ideally would yield one solution, but could result in many
 - ▶ I didn't see the problem space getting any smaller, and it was taking an hour+ for RI, which had 10 records
 - ▶ Generating list of permutations is time/memory expensive

Final Attempt

- Combination of previous attempts
- Generate one unique permutation/category at a time and assign to solution dataframe
- Check for consistency; if a category option can't be assigned, backtrack to previous category

Record ID	State	Event	Occupation	Industry	Public/private	Wage/self-employed	Gender	Age category	Race
1	42	1	40	230		2			
2	42	2	53	926	10				
3	42	5	11	561	50			5	6
4	42	4	47	230	50			6	
5	42	3	53	465					
6	42	3	47	713			2		
7	42	6	31	621					

Did it work?

Maybe!

- Produces a consistent table for small states relatively quickly; accuracy varies
- Can “stack” results to check for values that are the same between results and fix in place
- Can’t predict how many accurate solutions
- Can’t eliminate impossible permutation options
- Without original microdata to check, difficult to determine correctness
- Still slow for larger states



Results

	Solved	Acc.	Speed	Pros	Cons
Backtracking	X		X	<ul style="list-style-type: none"> • Easy to understand 	<ul style="list-style-type: none"> • Slow
OR-Tools Optimization	✓	X	✓	<ul style="list-style-type: none"> • Built-in functions 	<ul style="list-style-type: none"> • Inflexible
Merging Permutations	X		Okay	<ul style="list-style-type: none"> • Eliminate large numbers quickly 	<ul style="list-style-type: none"> • Slow overall • Uses lots of memory • Can't make per-value judgement
Combination	✓	Okay	✓	<ul style="list-style-type: none"> • Fastest • Multiple solutions • Look for consistency within variables 	<ul style="list-style-type: none"> • Can't eliminate possibilities that will never work

Going Forward

- Combine state results to check for consistency on a national level
- Test with data published following rule changes



Citations

- A. Narayanan and V. Shmatikov, "Robust De-anonymization of Large Sparse Datasets," *2008 IEEE Symposium on Security and Privacy (sp 2008)*, 2008, pp. 111-125, doi: 10.1109/SP.2008.33.
- S. Garfinkel, Abowd J.M, and Martindale C. 2018. "Understanding Reconstruction Attacks on Public Data," *Communications of the ACM*, 2018, vol. 62 no. 3, pp 46-43, doi: 10.1145/3287287
- [Google OR-Tools](#), Google Developers.

Citations



Contact Information

Sarah Birmingham

Coding It Forward Data Science Fellow
Census of Fatal Occupational Injuries

610-906-6593

sarahbirmingham42@gmail.com